

Data Profiling for the US Accidents (2016 - 2021) dataset

Doing some basic data profiling here using pandas - the main data profiling with a very elaborate data profiling report has been done as well and the report has been saved for cleansing usage

Also plotting a few graphs here using relevant libraries

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv('US_Accidents_Dec21_updated.csv')
#read the dataset
```

```
In [3]: df.columns
#printing all column names
```

```
Out[3]: Index(['ID', 'Severity', 'Start_Time', 'End_Time', 'Start_Lat', 'Start_Lng',
              'End_Lat', 'End_Lng', 'Distance(mi)', 'Description', 'Number', 'Street',
              'Side', 'City', 'County', 'State', 'Zipcode', 'Country', 'Timezone',
              'Airport_Code', 'Weather_Timestamp', 'Temperature(F)', 'Wind_Chill(F)',
              'Humidity(%)', 'Pressure(in)', 'Visibility(mi)', 'Wind_Direction',
              'Wind_Speed(mph)', 'Precipitation(in)', 'Weather_Condition', 'Amenity',
              'Bump', 'Crossing', 'Give_Way', 'Junction', 'No_Exit', 'Railway',
              'Roundabout', 'Station', 'Stop', 'Traffic_Calming', 'Traffic_Signal',
              'Turning_Loop', 'Sunrise_Sunset', 'Civil_Twilight', 'Nautical_Twilight',
              'Astronomical_Twilight'],
              dtype='object')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	ID	Severity	Start_Time	End_Time	Start_Lat	Start_Lng	End_Lat	End_Lng	Distance(mi)	Description	...	Roundabout	Station	Stop	Traffic_Calming
0	A-1	3	2016-02-08 00:37:08	2016-02-08 06:37:08	40.108910	-83.092860	40.112060	-83.031870	3.230	Between Sawmill Rd/Exit 20 and OH-315/Olentang...	...	False	False	False	Fal:
1	A-2	2	2016-02-08 05:56:20	2016-02-08 11:56:20	39.865420	-84.062800	39.865010	-84.048730	0.747	At OH-4/OH-235/Exit 41 - Accident.	...	False	False	False	Fal:
2	A-3	2	2016-02-08 06:15:39	2016-02-08 12:15:39	39.102660	-84.524680	39.102090	-84.523960	0.055	At I-71/US-50/Exit 1 - Accident.	...	False	False	False	Fal:
3	A-4	2	2016-02-08 06:51:45	2016-02-08 12:51:45	41.062130	-81.537840	41.062170	-81.535470	0.123	At Dart Ave/Exit 21 - Accident.	...	False	False	False	Fal:
4	A-5	3	2016-02-08 07:53:43	2016-02-08 13:53:43	39.172393	-84.492792	39.170476	-84.501798	0.500	At Mitchell Ave/Exit 6 - Accident.	...	False	False	False	Fal:

5 rows x 47 columns

```
In [5]: df.dtypes
#checking data types of all columns, this will help in deciding strategy for filling NULL/NaN values
```

```
Out[5]: ID                object
Severity                int64
Start_Time              object
End_Time                object
Start_Lat               float64
Start_Lng               float64
End_Lat                 float64
End_Lng                 float64
Distance(mi)            float64
Description              object
Number                  float64
Street                  object
Side                    object
City                    object
County                  object
State                   object
Zipcode                 object
Country                 object
Timezone                object
Airport_Code             object
Weather_Timestamp        object
Temperature(F)           float64
Wind_Chill(F)            float64
Humidity(%)              float64
Pressure(in)             float64
Visibility(mi)           float64
Wind_Direction           object
Wind_Speed(mph)          float64
Precipitation(in)        float64
Weather_Condition         object
Amenity                  bool
Bump                     bool
Crossing                 bool
Give_Way                 bool
Junction                 bool
No_Exit                  bool
Railway                  bool
Roundabout               bool
Station                  bool
Stop                     bool
Traffic_Calming           bool
Traffic_Signal            bool
Turning_Loop              bool
Sunrise_Sunset            object
Civil_Twilight            object
Nautical_Twilight         object
Astronomical_Twilight     object
dtype: object
```

```
In [6]: df.describe()
```


```
Out[6]:
```

	Severity	Start_Lat	Start_Lng	End_Lat	End_Lng	Distance(mi)	Number	Temperature(F)	Wind_Chill(F)	Humidity(%)	Pr
count	2.845342e+06	2.845342e+06	2.845342e+06	2.845342e+06	2.845342e+06	2.845342e+06	1.101431e+06	2.776068e+06	2.375699e+06	2.772250e+06	2.78
mean	2.137572e+00	3.624520e+01	-9.711463e+01	3.624532e+01	-9.711439e+01	7.026779e-01	8.089408e+03	6.179356e+01	5.965823e+01	6.436545e+01	2.94
std	4.787216e-01	5.363797e+00	1.831782e+01	5.363873e+00	1.831763e+01	1.560361e+00	1.836009e+04	1.862263e+01	2.116097e+01	2.287457e+01	1.04
min	1.000000e+00	2.456603e+01	-1.245481e+02	2.456601e+01	-1.245457e+02	0.000000e+00	0.000000e+00	-8.900000e+01	-8.900000e+01	1.000000e+00	0.00
25%	2.000000e+00	3.344517e+01	-1.180331e+02	3.344628e+01	-1.180333e+02	5.200000e-02	1.270000e+03	5.000000e+01	4.600000e+01	4.800000e+01	2.93
50%	2.000000e+00	3.609861e+01	-9.241808e+01	3.609799e+01	-9.241772e+01	2.440000e-01	4.007000e+03	6.400000e+01	6.300000e+01	6.700000e+01	2.98
75%	2.000000e+00	4.016024e+01	-8.037243e+01	4.016105e+01	-8.037338e+01	7.640000e-01	9.567000e+03	7.600000e+01	7.600000e+01	8.300000e+01	3.00
max	4.000000e+00	4.900058e+01	-6.711317e+01	4.907500e+01	-6.710924e+01	1.551860e+02	9.999997e+06	1.960000e+02	1.960000e+02	1.000000e+02	5.89

Calculate & visualize the percentage NULL values for each column or feature

```
In [7]: null_vals = df.isna().sum()/len(df)*100
df_null_vals = pd.DataFrame(null_vals)
df_null_vals.reset_index(inplace = True)
df_null_vals.columns = ["Feature", "Percent missing"]
plt.figure(figsize = (12,12))
plt.xticks(rotation = 45)
sns.barplot(x = "Percent missing", y = "Feature", data = df_null_vals, orient = "h")
```

```
Out[7]: <AxesSubplot:xlabel='Percent missing', ylabel='Feature'>
```



Print out the number of NaN/NULL values in each column

```
In [8]: df.isna().sum()
```

```
Out[8]: ID                0
Severity                0
Start_Time              0
End_Time                0
Start_Lat              0
Start_Lng              0
End_Lat                0
End_Lng                0
Distance(mi)           0
Description             0
Number                1743911
Street                  2
Side                    0
City                   137
County                 0
State                  0
Zipcode                1319
Country                0
Timezone               3659
Airport_Code           0
Weather_Timestamp      0
Temperature(F)         0
Wind_Chill(F)          17
Humidity(%)            3
Pressure(in)           3
Visibility(mi)          3
Wind_Direction         3
Wind_Speed(mph)        15
Precipitation(in)      19
Weather_Condition       3
Amenity                0
Bump                   0
Crossing               0
Give_Way               0
Junction               0
No_Exit                0
Railway                0
Roundabout             0
Station                0
Stop                   0
Traffic_Calming        0
Traffic_Signal         0
Turning_Loop           0
Sunrise_Sunset         0
Civil_Twilight         0
Nautical_Twilight      0
Astronomical_Twilight  0
```

```
In [9]: df_all_nulls = df[df.isna().any(axis=1)]
df_all_nulls
#printing all the rows which have a NULL value for any of the columns
```

Out[9]:

	ID	Severity	Start_Time	End_Time	Start_Lat	Start_Lng	End_Lat	End_Lng	Distance(mi)	Description	...	Roundabout	Station	Stop
0	A-1	3	2016-02-08 00:37:08	2016-02-08 06:37:08	40.108910	-83.092860	40.112060	-83.031870	3.230	Between Sawmill Rd/Exit 20 and OH-315/Olentang...	...	False	False	False
1	A-2	2	2016-02-08 05:56:20	2016-02-08 11:56:20	39.865420	-84.062800	39.865010	-84.048730	0.747	At OH-4/OH-235/Exit 41 - Accident.	...	False	False	False
2	A-3	2	2016-02-08 06:15:39	2016-02-08 12:15:39	39.102660	-84.524680	39.102090	-84.523960	0.055	At I-71/US-50/Exit 1 - Accident.	...	False	False	False
3	A-4	2	2016-02-08 06:51:45	2016-02-08 12:51:45	41.062130	-81.537840	41.062170	-81.535470	0.123	At Dart Ave/Exit 21 - Accident.	...	False	False	False
4	A-5	3	2016-02-08 07:53:43	2016-02-08 13:53:43	39.172393	-84.492792	39.170476	-84.501798	0.500	At Mitchell Ave/Exit 6 - Accident.	...	False	False	False
...
2845337	A-2845338	2	2019-08-23 18:03:25	2019-08-23 18:32:01	34.002480	-117.379360	33.998880	-117.370940	0.543	At Market St - Accident.	...	False	False	False
2845338	A-2845339	2	2019-08-23 19:11:30	2019-08-23 19:38:23	32.766960	-117.148060	32.765550	-117.153630	0.338	At Camino Del Rio/Mission Center Rd - Accident.	...	False	False	False
2845339	A-2845340	2	2019-08-23 19:00:21	2019-08-23 19:28:49	33.775450	-117.847790	33.777400	-117.857270	0.561	At Glassell St/Grand Ave - Accident. in the ri...	...	False	False	False
2845340	A-2845341	2	2019-08-23 19:00:21	2019-08-23 19:29:42	33.992460	-118.403020	33.983110	-118.395650	0.772	At CA-90/Marina Fwy/Jefferson Blvd - Accident.	...	False	False	False
2845341	A-2845342	2	2019-08-23 18:52:06	2019-08-23 19:21:31	34.133930	-117.230920	34.137360	-117.239340	0.537	At Highland Ave/Arden Ave - Accident.	...	False	False	False

1902024 rows x 47 columns

```
In [10]: df[df['Sunrise_Sunset'].isna()]
#printing rows which have NULL/NaN value for Sunrise_Sunset column
```


Out[10]:

	ID	Severity	Start_Time	End_Time	Start_Lat	Start_Lng	End_Lat	End_Lng	Distance(mi)	Description	...	Roundabout	Station	Stop	Traffic
34903	A-34904	4	2016-12-11 16:33:08	2016-12-11 22:33:08	41.03674	-73.67549	41.083450	-73.665300	3.271	Closed between CT-120A/King St/Exit 27 and Rou...	...	False	False	False	
43106	A-43107	4	2017-01-10 04:51:19	2017-01-10 10:51:19	40.84309	-84.80236	40.844250	-84.784350	0.945	Closed between State Line Rd and Kings Church	False	False	False	
43114	A-43115	4	2017-01-10 06:13:46	2017-01-10 12:13:46	40.84309	-84.80236	40.844250	-84.784350	0.945	Closed between State Line Rd and Kings Church	False	False	False	
										Closed between					

Visualizing the correlation between each column as a heatmap

```
In [11]: corr_matrix = df.corr()
plt.figure(figsize = (12,12))
sns.heatmap(corr_matrix, cmap = "YlGnBu")
```

Out[11]: <AxesSubplot:>



The following 3 cells are used to generate the data profiling report

The cells have been commented out since this has already ran before. The report is saved as a HTML document. In order to avoid running this again and using up a lot of time and memory, it has now been commented out

```
In [12]: #from pandas_profiling import ProfileReport
```

```
In [13]: #us_accidents_report = ProfileReport(df, title = "US Accidents Profile Report")
```

```
In [14]: #us_accidents_report.to_file("US_Accidents_Data_Profile_Group5.html")
```

Data Cleansing for the US Accidents (2016 - 2021) dataset

Changing the column name to exclude _ and make everything CamelCase

For columns which has units in them adding a _ to differentiate

```
In [15]: df.columns = df.columns.str.replace("[_,]", "", regex = True)
df.columns = df.columns.str.replace("(", "_", regex = False)
df.columns = df.columns.str.replace('Humidity_%', 'Humidity_Perc')
```

```
In [16]: df.columns
```

```
Out[16]: Index(['ID', 'Severity', 'StartTime', 'EndTime', 'StartLat', 'StartLng',
               'EndLat', 'EndLng', 'Distance_mi', 'Description', 'Number', 'Street',
               'Side', 'City', 'County', 'State', 'Zipcode', 'Country', 'Timezone',
               'AirportCode', 'WeatherTimestamp', 'Temperature_F', 'WindChill_F',
               'Humidity_Perc', 'Pressure_in', 'Visibility_mi', 'WindDirection',
               'WindSpeed_mph', 'Precipitation_in', 'WeatherCondition', 'Amenity',
               'Bump', 'Crossing', 'GiveWay', 'Junction', 'NoExit', 'Railway',
               'Roundabout', 'Station', 'Stop', 'TrafficCalming', 'TrafficSignal',
               'TurningLoop', 'SunriseSunset', 'CivilTwilight', 'NauticalTwilight',
               'AstronomicalTwilight'],
              dtype='object')
```

Change data type of timestamp columns to DateTime

```
In [17]: df[['StartTime', 'EndTime']] = df[['StartTime', 'EndTime']].apply(pd.to_datetime)
# doubt: can we change the column name/naming convention? Renamethis column to: StartTime and EndTime
```

The Weather_Timestamp column is also a timestamp column but it has NaN values. Replaces these NaN values with the maximum possible timestamp in Pandas, which is 2262-04-11 00:00:00

```
In [18]: df['WeatherTimestamp'] = df['WeatherTimestamp'].fillna('2262-04-11 00:00:00')
df['WeatherTimestamp'] = df['WeatherTimestamp'].apply(pd.to_datetime)
```

Creating a new 'AccidentYear' column which will be of help while doing visualization

```
In [19]: df['AccidentYear'] = df['StartTime'].dt.year
```

```
In [20]: df['AccidentYear'].value_counts()
```

```
Out[20]: 2021    1511745
         2020     625864
         2019     258615
         2017     163918
         2018     163176
         2016     122024
         Name: AccidentYear, dtype: int64
```

Filling all the NULL values in various column with placeholder respective placeholders

For all string columns, fill NULLs with 'Not listed'

```
In [21]: string_type_columns = ['Street', 'City', 'Zipcode', 'Timezone', 'AirportCode', 'WindDirection', 'WeatherCondition',
                                'SunriseSunset', 'CivilTwilight', 'NauticalTwilight', 'AstronomicalTwilight']
#first making a list of column names which are of string type and have NULL values
```

```
In [22]: df[string_type_columns] = df[string_type_columns].fillna('Not listed')
```

For all int/float columns, fill NULLs with -99

```
In [23]: integer_type_columns = ['Number', 'Temperature_F', 'WindChill_F', 'Humidity_Perc', 'Pressure_in', 'Visibility_mi',
                                'WindSpeed_mph', 'Precipitation_in']
```

```
In [24]: df[integer_type_columns] = df[integer_type_columns].fillna(-99)
```

No NULLs in the 'Side' column, but only one value is N. Replace it with the most frequent occurring value which is R

```
In [25]: df['Side'] = df['Side'].replace(['N'], 'R')
```

```
In [26]: df['Side'].value_counts() #just to verify that 'N' is replaced
```

```
Out[26]: R    2353310
         L     492032
         Name: Side, dtype: int64
```

```
In [27]: df.dtypes
```

```
Out[27]: ID                                object
Severity                                int64
StartTime                             datetime64[ns]
EndTime                             datetime64[ns]
StartLat                             float64
StartLng                             float64
EndLat                               float64
EndLng                               float64
Distance_mi                           float64
Description                           object
Number                               float64
Street                               object
Side                                 object
City                                 object
County                              object
State                               object
Zipcode                             object
Country                             object
Timezone                             object
..                                   ..
..                                   ..
```

```
In [28]: df.isna().sum()
```

```
Out[28]: ID                                0
Severity                                0
StartTime                             0
EndTime                             0
StartLat                             0
StartLng                             0
EndLat                               0
EndLng                               0
Distance_mi                           0
Description                           0
Number                               0
Street                               0
Side                                 0
City                                 0
County                              0
State                               0
Zipcode                             0
Country                             0
Timezone                             0
AirportCode                          0
WeatherTimestamp                     0
Temperature_F                        0
WindChill_F                          0
Humidity_Perc                        0
Pressure_in                          0
Visibility_mi                        0
WindDirection                        0
WindSpeed_mph                       0
Precipitation_in                     0
WeatherCondition                     0
Amenity                              0
Bump                                 0
Crossing                             0
GiveWay                              0
Junction                             0
NoExit                              0
Railway                              0
Roundabout                           0
Station                              0
Stop                                 0
TrafficCalming                       0
TrafficSignal                        0
TurningLoop                          0
SunriseSunset                        0
CivilTwilight                        0
NauticalTwilight                     0
AstronomicalTwilight                 0
AccidentYear                         0
dtype: int64
```

The following 2 new columns that are created in the 2 cells below are done for the data mapping purpose where we need a unique key value.

These 2 cells are only used for data loading into Neo4j and not important in terms of business

```
In [29]: #for Neo4j scripting purpose
df['TrafficCoordinates'] = df['StartLat'].astype(str) + ", " + df['StartLng'].astype(str) + ", " + df['EndLat'].astype
```

```
In [30]: df['TrafficDuration'] = df['StartTime'].astype(str) + ", " + df['EndTime'].astype(str)
```

No NaNs/NULLs remain in any column of the dataset. The dataset is cleaned!

```
In [31]: df.to_csv('US_Accidents_cleaned.csv', index=False)
#saving the cleaned data frame as a new CSV
```

```
In [2]: df = pd.read_csv('US_Accidents_cleaned.csv')
#read the dataset
```

```
In [3]: df
```

```
Out[3]:
```

	ID	Severity	StartTime	EndTime	StartLat	StartLng	EndLat	EndLng	Distance_mi	Description	...	TrafficCalming	TrafficSignal
0	A-1	3	2016-02-08 00:37:08	2016-02-08 06:37:08	40.108910	-83.092860	40.112060	-83.031870	3.230	Between Sawmill Rd/Exit 20 and OH-315/Olentang...	...	False	False
1	A-2	2	2016-02-08 05:56:20	2016-02-08 11:56:20	39.865420	-84.062800	39.865010	-84.048730	0.747	At OH-4/OH-235/Exit 41 - Accident.	...	False	False
2	A-3	2	2016-02-08 06:15:39	2016-02-08 12:15:39	39.102660	-84.524680	39.102090	-84.523960	0.055	At I-71/US-50/Exit 1 - Accident.	...	False	False
3	A-4	2	2016-02-08 06:51:45	2016-02-08 12:51:45	41.062130	-81.537840	41.062170	-81.535470	0.123	At Dart Ave/Exit 21 - Accident.	...	False	False
4	A-5	3	2016-02-08 07:53:43	2016-02-08 13:53:43	39.172393	-84.492792	39.170476	-84.501798	0.500	At Mitchell Ave/Exit 6 - Accident.	...	False	False
...
2845337	A-2845338	2	2019-08-23 18:03:25	2019-08-23 18:32:01	34.002480	-117.379360	33.998880	-117.370940	0.543	At Market St - Accident.	...	False	False
2845338	A-2845339	2	2019-08-23 19:11:30	2019-08-23 19:38:23	32.766960	-117.148060	32.765550	-117.153630	0.338	At Camino Del Rio/Mission Center Rd - Accident.	...	False	False
2845339	A-2845340	2	2019-08-23 19:00:21	2019-08-23 19:28:49	33.775450	-117.847790	33.777400	-117.857270	0.561	At Glassell St/Grand Ave - Accident. in the ri...	...	False	False
2845340	A-2845341	2	2019-08-23 19:00:21	2019-08-23 19:29:42	33.992460	-118.403020	33.983110	-118.395650	0.772	At CA-90/Marina Fwy/Jefferson Blvd - Accident.	...	False	False
2845341	A-2845342	2	2019-08-23 18:52:06	2019-08-23 19:21:31	34.133930	-117.230920	34.137360	-117.239340	0.537	At Highland Ave/Arden Ave - Accident.	...	False	False

2845342 rows × 50 columns

```
In [4]: df['AccidentYear'].value_counts()
```

```
Out[4]:
```

2021	1511745
2020	625864
2019	258615
2017	163918
2018	163176
2016	122024

Name: AccidentYear, dtype: int64

```
In [7]: df_sampled = df.sample(frac = .01, ignore_index = True)
```


In [8]: df_sampled

Out[8]:

	ID	Severity	StartTime	EndTime	StartLat	StartLng	EndLat	EndLng	Distance_mi	Description	...	TrafficCalming
0	A-408764	2	2021-11-18 01:07:00.000000000	2021-11-18 02:22:09.000000000	30.442862	-85.874004	30.443136	-85.874010	0.019	Incident on DOG TRACK RD near CAPTAIN FRITZ RD...	...	False
1	A-645069	2	2021-11-12 06:14:30	2021-11-12 07:33:18	38.269060	-77.945498	38.269320	-77.928368	0.929	On VA-20 in the County of Orange in the vicini...	...	False
2	A-600734	2	2021-09-23 16:37:31	2021-09-23 17:54:19	34.284416	-79.909065	34.285379	-79.910967	0.127	Incident on ROGERS RD near INDIAN BRANCH RD Ex...	...	False
3	A-1950772	2	2020-12-07 12:23:30	2020-12-08 04:06:01	28.428795	-81.407121	28.428975	-81.404630	0.152	Slow traffic on US-92 E - US-17 N - US-441 N f...	...	False
4	A-388339	2	2021-11-09 18:59:00	2021-11-09 20:22:38	33.638815	-117.344412	33.638426	-117.343770	0.046	Incident on GRAND AVE near HOUSE 17671 Right S...	...	False
...
28448	A-109926	2	2017-01-10 13:53:17	2017-01-10 19:53:17	37.560820	-122.302940	37.563630	-122.306010	0.257	At Kehoe Ave/Exit 415 - Accident.	...	False
28449	A-1344523	2	2021-06-28 19:55:00	2021-06-28 20:52:30	34.140763	-117.248329	34.135979	-117.236718	0.742	Slow traffic on CA-210 E from Del Rosa Ave (CA...	...	False
28450	A-2419578	2	2019-10-16 10:58:00	2019-10-16 12:15:00	40.388240	-111.835720	40.388240	-111.835380	0.018	At I-15 - Accident.	...	False
28451	A-971142	2	2021-10-26 15:31:04	2021-10-26 15:51:31	25.857208	-80.209463	25.861819	-80.209631	0.319	Stationary traffic on US-441 N from FL-934/81s...	...	False
28452	A-743597	2	2021-12-08 15:10:30	2021-12-08 19:17:00	33.715089	-117.845076	33.722496	-117.837676	0.665	Slow traffic on CA-55 N - Costa Mesa Fwy N fro...	...	False

28453 rows x 50 columns

In [9]: df_sampled['AccidentYear'].value_counts()

Out[9]:

2021	15106
2020	6289
2019	2624
2018	1705
2017	1602
2016	1127

Name: AccidentYear, dtype: int64

In [10]: df_sampled.to_csv('US_Accidents_sampled.csv', index=False)
#saving the cleaned data frame as a new CSV

In []: